Performance Impact of Undefined Behavior Optimizations in C/C++

Lucian Popescu*, Razvan Deaconescu* and Nuno Lopes**

*Facultatea de Automatică și Calculatoare, Universitatea Politehnică din București
**Instituto Superior Técnico, Universidade de Lisboa
*lucian.popescu187@gmail.com, razvan.deaconescu@cs.pub.ro
**nuno.lopes@tecnico.ulisboa.pt

Wednesday 14th June, 2023

Abstract

Clang/LLVM uses undefined behavior to issue optimizations. In this report, we present the impact of this class of optimizations on a benchmarking suite built using Phoronix. In this suite we focus on a diverse set of application categories, ranging from compression algorithms and image processing to web-servers and databases. To extensively cover the impact on these applications, we provide 10 configurations that modify the behavior of the compiler when exploiting undefined behavior in optimizations. Current results show that in 90% of cases the impact of undefined behavior optimizations is insignificant.

1 Introduction

Our work is focused on evaluating the performance impact of undefined behavior optimizations. Traditionally, undefined behavior was used as a free ticket to issue various types of optimizations targeting either code speed or code size. The impact of this type of optimization is already well known for micro-benchmarks, however, there is no study covering the impact for real loads. We fill this gap by benchmarking various instances of undefined behavior optimizations in Clang/LLVM on real application loads written in C/C++.

The work in this semester was broken down into two parts. First, we started by creating the infrastructure for running the benchmarks. Our main goal in this regard was to gather a set of popular applications written in C and C++. After exploring existing solutions for this problem, we decided to use the Phoronix Test Suite for this task. The applications provided by Phoronix were later compiled with various configurations that modify the behavior of the compiler when exploiting undefined behavior. For each configuration, we gathered a set of results that we will present later in this work.

Second, we were interested in exploring configurations for undefined behavior exploitation. We started from a set of already-implemented configurations, such as -fwrapv or -fno-

strict-aliasing, and moved to configurations implemented by us, such as -fno-constrain-bool-value. We benefited from 5 already-implemented configurations and we managed to implement 5 more configurations. For some of the configurations the level of difficulty was relatively low, because we only had to modify the frontend of the compiler, i.e. Clang. However there existed cases where we also had to modify the middle-end of the compiler, i.e. LLVM.

This work is structured as follows: Section 2 provides detailed information about the process of creating the benchmarks for our use case, Section 3 talks about each compiler configuration that modifies the behavior of the compiler when exploiting undefined behavior, Section 4 presents the results of the benchmarks for each configuration, and Section 5 discusses the limitations of the benchmarking process.

2   Creating the Benchmarking Infrastructure

To simplify the benchmarking process, we needed an infrastructure that would be triggered for each undefined behavior configuration. The requirements for such an infrastructure were: contain benchmarks for C/C++ applications and avoid synthetic benchmarks. We wanted to avoid synthetic benchmarks because we aimed to evaluate the impact of undefined behavior optimizations on real application loads.

The investigation of the state-of-the-art benchmarking infrastructure for real application loads resulted in two candidates, i.e. Phoronix Test Suite [5] and Geekbench [2]. We chose to go further with Phoronix because its software is open-source and it can be extended with new benchmarks as per our needs, as opposed to Geekbench.

Phoronix offers a wide variety of benchmarks, written in various programming languages. We were only interested in the ones written in C and C++. To do that we had to filter all benchmarks by their dependencies and choose only the ones that had a C/C++ compiler as a dependency. This resulted in a list of nearly 200 applications.

We had to further filter out the benchmark applications because a number of them had problems while they were compiled with Clang. The effort of making them work with this compiler was not worth as we had a good amount of applications that already worked.

Next, we had the problem of benchmarking time. Spending too much time on benchmarking was not beneficial for us so we decided to limit one round of benchmarking to 24 hours. Because of that we had to further cut applications that required a significant amount of time to benchmark. One example of such application is GCC. One of the benchmarks in our suite was measuring the time of building GCC form sources. GCC was a complicated 3-step build process, that takes a few hours to complete, thus we wanted to avoid it.

This process resulted in the benchmark suite presented in Table 1. At this moment the suite mostly contains applications that are either CPU bound or memory bound. We discarded completely GPU, OpenMP and MPI applications as they are more difficult to benchmark.

3 Compiler Configurations based on Undefined Behavior

After we defined the benchmark suite, we focused on evaluating the performance impact of various undefined behavior optimizations. In this section we present first the already-implemented configurations for the above mentioned optimizations and second we present the configurations that we implemented.

Each configuration is made available through a compiler flag that can be used when compiling the benchmarks.

3.1 Already-implemented compiler flags

We benefited from 5 flags in this category.

-fwrapv instructs the compiler to assume that signed arithmetic overflow of addition, subtraction and multiplication wraps around using twos-complement representation. In LLVM, this has the impact of dropping the *nsw* [4] attribute in the above mentioned arithmetic operations.

-fno-strict-aliasing instructs the compiler to apply the strictest aliasing rules available. In LLVM, this has the impact of dropping the *tbaa* [7] attribute that is used for type based alias analysis.

-fstrict-enums instructs the compiler to optimize using the assumption that a value of enumerated type can only be one of the values of the enumeration (as defined in the C++ standard; basically, a value that can be represented in the minimum number of bits needed to represent all the enumerators). In LLVM, this has the impact of adding the *range* [6] attribute to memory operations.

-fno-delete-null-pointer-checks instructs the compiler to assume that programs can safely dereference NULL pointers and thus to not delete NULL pointer checks that are proved to be redundant.

-fno-finite-loops instructs the compiler to assume that no loop is finite. In LLVM, this has the impact of dropping the *mustprogress* [3] attribute from all loops and functions.

3.2 Added compiler flags

-fconstrain-shift-value instructs the compiler to mask the right-hand-side (RHS) of the shift operation so that it does not produce undefined behavior when the RHS is bigger that the bitwidth. On x86, this add an additional *and* instruction for masking.

-fno-constrain-bool-value instructs the compiler to not constrain bool values to 0 and 1. In LLVM, this has the impact of dropping the *range* attribute from memory operations that work with booleans.

-fno-use-default-alignment instructs the compiler to use alignment 1 for all memory operations including load, store, memcpy, etc. The alignments of global variables and allocas remain unaffected. This has the impact of not generating the most efficient code because the compiler cannot find the best alignment for each operation that was forcefully aligned to 1.

-mllvm -zero-uninit-loads instructs the compiler to replace uninitialized loads with zero loads. This does not automatically initialize all memory with zero, instead it fills the memory with zero only when the memory is requested.

-fdrop-inbounds-from-gep -mllvm -trap-on-oob instructs the compiler to trap when it detects an out-of-bounds (OOB) memory access. By trapping, the compiler is blocked from doing any further optimizations based on OOB. This is a combination between a Clang flag (*-fdrop-inbounds-from-gep*) and a LLVM flag(*-trap-on-oob*). We needed the Clang flag to make sure that no optimization is triggered on the OOB access before we add the trap.

At the moment of writing this report, the last two flags are still in development.

3.3 Implementation details for added compiler flags

3.3.1 -fconstrain-shift-value

To implement *-fconstrain-shift-value* we had to do modifications in the frontend of the compiler. The OpenCL specification already implemented this behavior in EmitShl and EmitShr functions. Thus, we extended this by adding the new flag for our purposes, as presented in Listing 1.

```
1 @@ −3936,7 +3936,7 @@ Value *ScalarExprEmitter::EmitShl(const BinOpInfo &
      Ops) {
2    bool SanitizeBase = SanitizeSignedBase || SanitizeUnsignedBase;
3    bool SanitizeExponent = CGF.SanOpts.has(SanitizerKind::ShiftExponent);
4    // OpenCL 6.3j: shift values are effectively % word size of LHS.
5 −  if (CGF.getLangOpts().OpenCL)
6 +  if (CGF.getLangOpts().OpenCL || CGF.CGM.getCodeGenOpts().
      ConstrainShiftValue)
7      RHS = ConstrainShiftValue(Ops.LHS, RHS, "shl.mask");
8    else if ((SanitizeBase || SanitizeExponent) &&
9            isa<llvm::IntegerType>(Ops.LHS−>getType())) {
10 @@ −4005,7 +4005,7 @@ Value *ScalarExprEmitter::EmitShr(const BinOpInfo &
      Ops) {
11      RHS = Builder.CreateIntCast(RHS, Ops.LHS−>getType(), false, "sh_prom")
      ;
12
13    // OpenCL 6.3j: shift values are effectively % word size of LHS.
14 −  if (CGF.getLangOpts().OpenCL)
15 +  if (CGF.getLangOpts().OpenCL || CGF.CGM.getCodeGenOpts().
      ConstrainShiftValue)
16      RHS = ConstrainShiftValue(Ops.LHS, RHS, "shr.mask");
17    else if (CGF.SanOpts.has(SanitizerKind::ShiftExponent) &&
18            isa<llvm::IntegerType>(Ops.LHS−>getType())) {
```

Listing 1: -fconstrain-shift-value implementation

The resulted LLVM IR for this change is presented in Listing 2. For the constrained version, there is an additional and instruction introduced for clamping the RHS of the shift operation.

```
1 diff oversized−shift−constrain.ll oversized−shift−no−constrain.ll
2 <   %shl.mask = and i32 %1, 31
3 <   %shl = shl i32 %0, %shl.mask
4 −−−
```

```
5 >    %shl = shl i32 %0, %1
```

<div align="center">Listing 2: -fconstrain-shift-value results</div>

3.3.2 -fno-constrain-bool-value

For *-fno-constrain-bool-value* we used the same strategy as above, i.e. modify the frontend of the compiler to propagate the changes down to LLVM's Intermediate Representation (IR), as presented in Listing 3.

```
1 @@ −1683,7 +1683,7 @@ static bool getRangeForType(CodeGenFunction &CGF,
      QualType Ty,
2  llvm::MDNode ∗CodeGenFunction::getRangeForLoadFromType(QualType Ty) {
3    llvm::APInt Min, End;
4    if (!getRangeForType(∗this, Ty, Min, End, CGM.getCodeGenOpts().
     StrictEnums,
5 −                     hasBooleanRepresentation(Ty)))
6 +                     hasBooleanRepresentation(Ty) && CGM.getCodeGenOpts
     ().ConstrainBoolValue))
7      return nullptr;
8
9    llvm::MDBuilder MDHelper(getLLVMContext());
```

<div align="center">Listing 3: -fnoconstrain-bool-value implementation</div>

The resulted LLVM IR for the above change will be similar to the one presented in Listing 4. For the constrained (default) version, the !9 attribute is used for constraining the bool value between 0 and 2. For the unconstrained version, the !9 attribute is dropped and an additional and instruction is added in order to preserve the constraints of the type. However because this lies in the domain of undefined behavior, the and instruction is not required to be present.

```
1 −−− constrain−bool.ll
2 +++ no−constrain−bool.ll
3  define dso_local noundef zeroext i1 @_Z3fooPb(ptr nocapture noundef
     readonly
4 %a) local_unnamed_addr #0 {
5  entry:
6 −   %0 = load i8, ptr %a, align 1, !tbaa !5, !range !9
7 −   %tobool = icmp ne i8 %0, 0
8 +   %0 = load i8, ptr %a, align 1, !tbaa !5
9 +   %1 = and i8 %0, 1
10 +   %tobool = icmp ne i8 %1, 0
11    ret i1 %tobool
12  }
13
14 @@ −25,4 +26,3 @@
15  !6 = !{!"bool", !7, i64 0}
16  !7 = !{!"omnipotent char", !8, i64 0}
17  !8 = !{!"Simple C++ TBAA"}
18 −!9 = !{i8 0, i8 2}
```

<div align="center">Listing 4: -fnoconstrain-bool-value results</div>

3.3.3 -fno-use-default-alignment

To implement this flag we had to modify many parts of Clang. Because we are working with memory operations in this flag, we had to patch all calls to CreateStore, CreateLoad,

CreateMemCpy, etc [1] methods. This was a non-trivial task as there are tens of calls to these methods in different subsystems of Clang.

After patching all usages, we discovered the problem of atomic loads and stores. On x86, these types of memory operations need to be properly aligned, otherwise the processor will trap. To solve this problem, we decided to link the compiled applications with the atomic library provided by GCC which takes care of atomic memory operations.

3.3.4   -zero-uninit-loads and -trap-on-oob

For implementing these flags we had do to modifications in the middle-end of the compiler. They are still in progress as the level of difficulty for implementing them is greater than the previous flags.

For *-zero-uninit-loads* the implementation difficulty is to find the places where various types of variables get their initial values and replace uninitialized values for each type with zero values for the corresponding types.

For *-trap-on-oob* we plan to add traps for out-of-bounds accesses and stop further optimizations from taking place. The difficulty for this task is to spot all the places where the compiler exploits out-of-bounds accesses as there are many places where this happens.

4   Results

The benchmarks were run on a machine with the following specifications:

- Processor: 2 x Intel Xeon E5-2680 v2 @ 3.60GHz (20 Cores / 40 Threads), MicroArch: IvyBridge

- OS: Debian 11, kernel: 5.10.0-21-amd64 (x86_64)

- Compiler: Clang 15.0.7

To gather relevant results, we compiled the whole benchmark suite presented in Table 1 with a single flag presented in Section 3 at a time. Furthermore we used as baseline the benchmarks compiled with no flag presented in Section 3. After this step we used Phoronix to run the generated binaries for each benchmark.

While running a benchmark, Phoronix tries to reduce the noise as much as possible by rerunning the benchmark until the deviation between the results becomes minimal. This helped us because we did not have to do any further calibration to the results.

After compiling and running the benchmarks with all the configurations, we started to compare each configuration against the baseline. In this step we recorded the percentage of positive and negative performance impact relative to the baseline. To take into account noise in the results, we defined the noise threshold as -2% for negative performance impact and +2% for positive performance impact.

Plots presenting the performance impact for each flag are presented in Subsection 8.1 from Appendix.

For each flag, in nearly 90% of the cases, the performance impact is insignificant, i.e. can be considered noise as it is placed between -2% and +2%.

-fwrapv has the biggest overall negative performance impact, i.e. 11% of the results for *-fwrapv* have negative performance impact. This flag also exhibits the highest negative performance impact as presented in Figure 1. Other benchmarks with negative impact: FFTW - Float + SSE - Size: 1D FFT Size 256, uvg266 (Video Encoder). Other benchmarks with positive impact: OpenSSL - RSA4096.

-fno-strict-aliasing is the most balanced flag up until this moment. 3% of the results are positive performance impact and 8.2% of the are negative performance impact. The outliers for this flag are also very close to the noise thresholds.

-fconstrain-shift-value has many results in the positive impact area, i.e. 8.1%. It also has the biggest overall positive performance impact, i.e. it increases the performance with 35% relative to baseline. This outlier exhibits for the benchmark presented in Figure 2. In this benchmark, *-fconstrain-bool-value* also presents a considerable positive performance impact.

-fno-use-default-alignment exhibits 2 times more positive performance impact than negative performance impact. This result is unexpected because in theory the behavior that this flag exhibits should decrease the performance of memory accesses.


Figure 1: eSpeak-NG Speech Engine - Text-To-Speech Synthesis Benchmark, Baseline: 41.59 Seconds

Figure 2: GtkPerf - GTK Widget: GtkDrawingArea - PixBufs, Baseline: 170.08 Seconds

5   Discussion

In this section we discuss current limitations regarding the benchmarking process for evaluating the performance impact of undefined behavior optimizations.

5.1   Categorization of Benchmarks

Current benchmark suite needs more investigation so that we understand what are its biases and its limitations. Our initial strategy of creating this benchmark was to gather as much applications written in C/C++ and that contain real loads, as opposed to synthetic loads. However we did no categorization to understand whether we cover a wide range of application or whether the current application categories overlap.

To do this categorization work, we will have to go through each benchmark and discover language patterns that apply to more benchmarks. Such a language pattern might be the usage of hot loops that contain most of the computation and logic of the program. Another pattern might be the heavy usage of network communication.

After we finish this categorization task, we can understand better the results and add new categories as needed.

5.2   Benchmarking Strategy

To gather meaningful results for our benchmarks we rely on Phoronix's feature of rerunning each benchmark a number of times so that the standard error between the data points is at a minimum. However this might give unreliable results as in the general case the distribution of the results is not standard. To fill this gap we can use another feature of Phoronix that returns the confidence interval for each result.

Using the above mentioned features together, we can spot the benchmarks for which a higher degree of noise was introduced and try to reduce it.

5.3 Discovering New Undefined Behavior Optimizations

Even if all undefined behavior instances are described in the C or C++ languages standard, the compiler might not make use of all of them. Up until this moment we relied on our experience to discover new undefined behavior optimizations. However, in the future we want to rely on automated tools to spot the optimization passes that exploit undefined behavior. To do that we plan to use Alive2 [8], which is an automatic verification tool for LLVM optimization passes.

6   Conclusions

In this semester we started the work of evaluating the performance impact of undefined behavior optimizations. We split the work in two parts. In the first part we focused on developing a benchmark suite that contains C/C++ applications with real loads, as opposed to synthetic loads. Then we started benchmarking 10 compiler flags that control the behavior of the compiler optimizations with regards to undefined behavior. 5 of them were already implemented, but we also added 5 new flags. Early results show that in nearly 90% of the cases the performance impact of undefined behavior in optimizations is insignificant.

7   Further Work

In the next semester we plan to explore new compiler configurations. This includes optimizations based on use-after-free, optimizations based on alias analysis that uses object-based rules or optimizations based on arithmetic related undefined behaviors, such as division by 0.

We also plan to run the benchmarks on other hardware architectures such as AMD or ARM to explore how they behave in comparison with the current hardware setup that is based on Intel.

It's also an interest for us to combine the flags to see how they behave together with regards to performance but must important is to analyze the current impact and discover the root causes of current performance numbers.

References

[1] clang::CodeGen::CGBuilderTy Class Reference. `https://clang.llvm.org/doxygen/classclang_1_1CodeGen_1_1CGBuilderTy.html`, last visited Wednesday 14th June, 2023.

[2] Geekbench. `https://www.geekbench.com/`, last visited Wednesday 14th June, 2023.

[3] mustprogress definition in LangRef. `https://llvm.org/docs/LangRef.html#llvm-loop-mustprogress-metadata`, last visited Wednesday 14th June, 2023.

[4] nsw definition in LangRef. `https://llvm.org/docs/LangRef.html#id89`, last visited Wednesday 14th June, 2023.

[5] Phoronix Test Suite. `https://www.phoronix-test-suite.com/`, last visited Wednesday 14th June, 2023.

[6] range definition in LangRef. `https://llvm.org/docs/LangRef.html#range-metadata`, last visited Wednesday 14th June, 2023.

[7] tbaa definition in LangRef. `https://llvm.org/docs/LangRef.html#tbaa-metadata`, last visited Wednesday 14th June, 2023.

[8] N. P. Lopes, J. Lee, C.-K. Hur, Z. Liu, and J. Regehr. Alive2: bounded translation validation for llvm. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, pages 65–79, 2021.

8 Appendix

Benchmark Suite	
Application Category	Phoronix Application Identifier
LLVM Build Speed	build-llvm
HPC	fftw
Video Encoding	aom-av1
	uvg266
Simulation	brl-cad
Bioinformatics	mrbayes
	hmmer
Image Processsing	jpegxl
	graphics-magick
Raytracing	tungsten
Parallel Processing	tjbench
	simdjson
Security	aircrack-ng
	openssl
Password Cracking	john-the-ripper
Database	redis
Audio Encoding	encode-flac
Texture Compression	basis
	draco
Compression	compress-zstd
	compress-pbzip2
Speech	espeak
	rnnoise
Software Defined Radio	liquid-dsp
GUI	gtkperf
Finance	quantlib
Telephony	pjsip
Circuit Simulator	ngspice
Webserver	apache
	nginx
Theorem Prover	z3

Table 1: Final benchmark suite that uses Phoronix applications. Second column contains Phoronix application identifiers, more information about them can be found at https://openbenchmarking.org/

8.1 Performance Impact for Undefined Behavior Flags
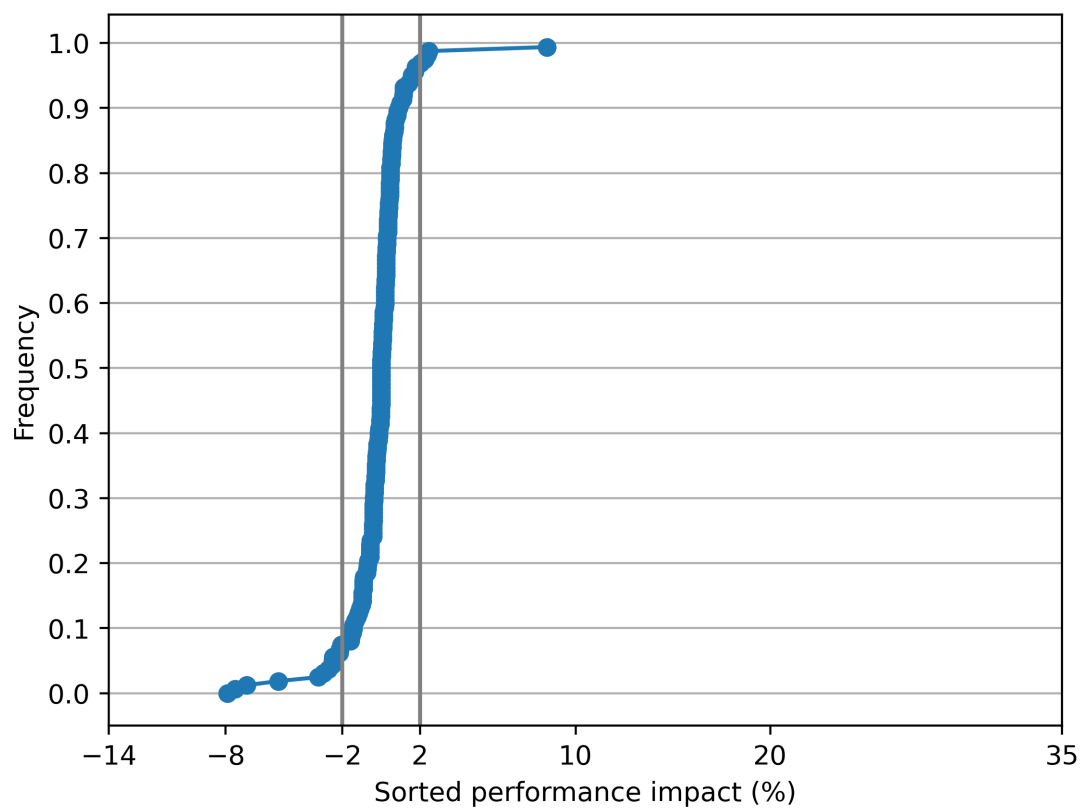
Figure 3: CDF of performance impact for -fwrapv

Figure 4: CDF of performance impact for -fno-strict-aliasing

Figure 5: CDF of performance impact for -fstrict-enums

Figure 6: CDF of performance impact for -fno-delete-null-pointer-checks

Figure 7: CDF of performance impact for -fconstrain-shift-value

Figure 8: CDF of performance impact for -fno-constrain-bool-value

Figure 9: CDF of performance impact for -fno-use-default-alignment